

Applied Machine Learning

Kernel Methods


Computer Science, Fall 2022

Instructor: Xuhong Zhang

Kernel Method

- Not matter hard SVM or soft SVM, we assume linear separation (w/o compromise).
- **Question:** What if the surface is **Non-Linear**?

Kernel Method



1, Can we use a linear classifier to classify the data?

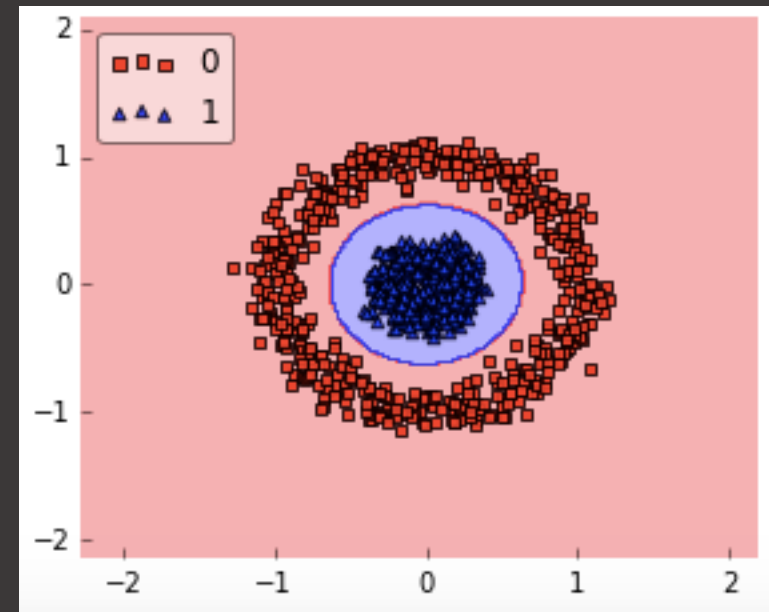

2, Can I fine-tune a linear classifier to classify the data?

3, What if I just want to use a linear classifier?

Kernel Method

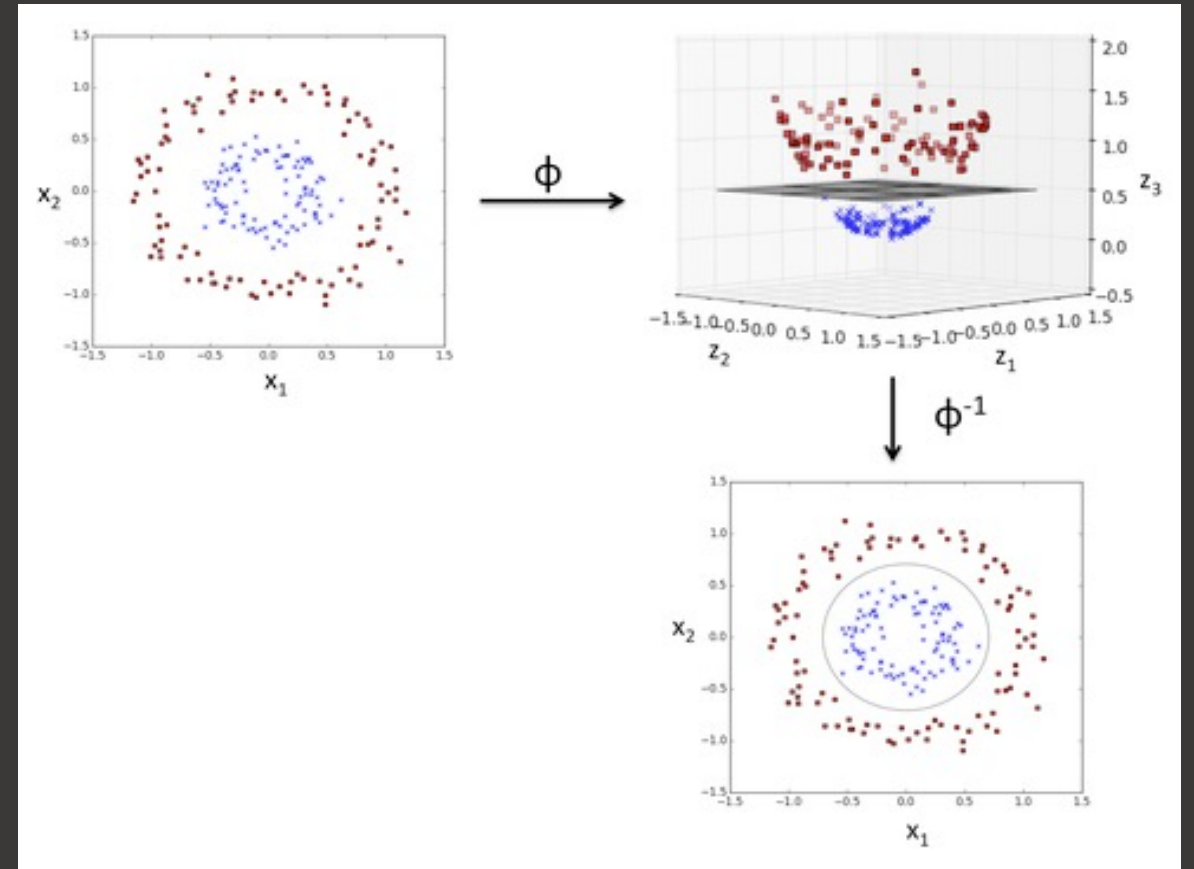

- The motivation for kernel method:

Making the Non-Linear Linear



Kernel Method

The RBF Kernel



Question: In order to get this decision boundary, what new features I can create to make this data separatable by this circular decision boundary?

Kernel Method

<https://www.youtube.com/watch?v=3liCbRZPrZA>

Kernel Method


- Motivation: we want to add/create new features so that the data is separatable by trivial linear classifier.
- Kernel method:

Making the Non-Linear Linear

Kernel Method

- **The key** : Mapping into a new Feature Space

$$\Phi: \chi \mapsto \hat{\chi} = \Phi(x)$$



Input Space

Feature Space

Recall what we talked about
high dimensional space

- For example, with $x_i \in \mathbb{R}^2$
We can construct

$$\Phi([x_{i1}, x_{i2}]) = [x_{i1}, x_{i2}, x_{i1}x_{i2}, x_{i1}^2, x_{i2}^2]$$

Handcrafted Feature Expansion

- By applying basis function (feature transformations) on the input feature vectors, we can make linear classifiers non-linear.
- For a data vector $x \in \mathbb{R}^d$, we apply the transformation $x \rightarrow \phi(x)$ where $\phi(x) \in \mathbb{R}^D$ (usually $D \gg d$)
- New question: not every problem is like the one we just saw and we don't know what features we need to add/create....

Handcrafted Feature Expansion

- Solution: I just add all the features we can think of...
- Example:

$$x = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{pmatrix}$$

$$\phi(x) = \begin{pmatrix} x_1 \\ \vdots \\ x_d \\ x_1 x_2 \\ \vdots \\ x_{d-1} x_d \\ \vdots \\ x_1 x_2 \dots x_d \end{pmatrix}$$


$\phi(x)$ is very expressive and allows for complicated non-linear decision boundaries, but the dimensionality is extremely high.

Question: what is the dimension of $\phi(x)$

Kernel Method

- In practice, we will apply two tricks to apply kernel methods.
 - The way to add all the interaction terms will cause memory issue and also prohibitively expensive for computation.
 - First trick: We have a linear classifier, and it turns out to be, we can express the linear classifier with the expression of inner product, e.g., $K_{ij} = x_i^T x_j$. We will take the square loss as an example.
 - If we only access our data with the inner product, we can pre-compute and store them.

Gradient Descent for Squared Loss

- For a linear model with the format of $w^T x_j$, claim that the w can be expressed as a linear combination of x , $\sum_{i=1}^n \alpha_i x_i$.
- If so, then the format $w^T x_j = \sum_{i=1}^n \alpha_i \boxed{x_i^T x_j}$  K_{ij}
- If everything can be expressed in the inner-product, then the inner-product can be pre-computed and stored. Save computational time. No matter how high dimension the data is, it is tolerable.

Gradient Descent for Squared Loss

Some observation:

- Look at the squared loss:

$$l(w) = \sum_{i=1}^n (w^T x_i - y_i)^2$$

- The gradient descent rule, with step-size/learning-rate $s > 0$, w is updated as:

$$w_{t+1} \leftarrow w_t - s \left(\frac{\partial l}{\partial w} \right), \text{ where } \frac{\partial l}{\partial w} = \sum_{i=1}^n 2(w^T x_i - y_i) x_i = \sum_{i=1}^n \gamma_i x_i$$

the gradient is a linear combination of input samples.

γ_i : function of x_i, y_i
And it is a scalar.

Proof

- Claim: w is a linear combination of all input vectors.

$$w = \sum_{i=1}^n \alpha_i x_i$$

- Proof :

1. Since the loss is convex, the final solution is independent of the initialization, and we can initialize w^0 to be whatever we want. So we pick $w^0 = [0, \dots, 0]^T$ and here we have $\alpha_1 = \dots = \alpha_n = 0$.

Proof

- Proof :

2. We re-write the gradient updates entirely in terms of updating the α_i coefficients

$$w_1 = w_0 - s \sum_{i=1}^n 2(w_0^T x_i - y_i)x_i = \sum_{i=1}^n \alpha_i^0 x_i - s \sum_{i=1}^n \gamma_i^0 x_i = \sum_{i=1}^n \alpha_i^1 x_i \quad (\text{with } \alpha_i^1 = \alpha_i^0 - s\gamma_i^0)$$

$$w_2 = w_1 - s \sum_{i=1}^n 2(w_1^T x_i - y_i)x_i = \sum_{i=1}^n \alpha_i^1 x_i - s \sum_{i=1}^n \gamma_i^1 x_i = \sum_{i=1}^n \alpha_i^2 x_i \quad (\text{with } \alpha_i^2 = \alpha_i^1 - s\gamma_i^1)$$

$$w_3 = w_2 - s \sum_{i=1}^n 2(w_2^T x_i - y_i)x_i = \sum_{i=1}^n \alpha_i^2 x_i - s \sum_{i=1}^n \gamma_i^2 x_i = \sum_{i=1}^n \alpha_i^3 x_i \quad (\text{with } \alpha_i^3 = \alpha_i^2 - s\gamma_i^2)$$

.....

Induction proof

$$w_t = w_{t-1} - s \sum_{i=1}^n 2(w_{t-1}^T x_i - y_i)x_i = \boxed{\sum_{i=1}^n \alpha_i^{t-1} x_i - s \sum_{i=1}^n \gamma_i^{t-1} x_i} = \sum_{i=1}^n \alpha_i^t x_i \quad (\text{with } \alpha_i^t = \alpha_i^{t-1} - s\gamma_i^{t-1})$$

Proof

- Put everything together:
 1. w is trivially a linear combination of the training vectors for w_0 (base case).
 2. If we apply the inductive hypothesis for w_t it follows for w_{t+1} .
 3. The update-rule for α_i^t is thus

$$\alpha_i^t = \alpha_i^{t-1} - s\gamma_i^{t-1}, \text{ and we have } \alpha_i^t = -s \sum_{r=0}^{t-1} \gamma_i^r$$

We only need to keep track of the n coefficients $\alpha_1, \dots, \alpha_n$, and they are independent of the dimensionality of the data ; only scale with n ; no need to compute w !

Good news: we can perform the entire gradient descent update rule without ever expressing w explicitly.

Loss function

In addition, we can also express the inner-product of w with any input x_i purely in terms of inner-products between training inputs.

$$w^T x_j = \sum_{i=1}^n \alpha_i x_i^T x_j$$

Thus, if we come back to the squared loss function, $l(w)$, we can rewrite it in terms of inner-product between training inputs:

$$l(w) = \sum_{i=1}^n (w^T x_i - y_i)^2 \quad \longleftrightarrow \quad l(\alpha) = \sum_{i=1}^n \left(\sum_{j=1}^n \alpha_j x_j^T x_i - y_i \right)^2$$

Benefit

During the test-time we also only need these coefficients to make a prediction on a test-point x_t , and can write the entire classifier in terms of inner-products between the test point and training points:

$$h(x_t) = w^T x_t = \sum_{i=1}^n \alpha_i x_i^T x_t$$

The only information we ever need in order to learn a hyper-plane classifier with the squared-loss is inner-products between all pairs of data vectors.

Inner-Product

- Now we get the idea that we can express our linear model using inner products. Remember this is the first trick.
- Remember that we want to add/create new features for the model. With the inner-product, $x_i^T x_j$, we have $\phi(x_i)^T \phi(x_j)$, and we just adopt

$$\phi(x) = \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_d \\ x_1 x_2 \\ \vdots \\ x_{d-1} x_d \\ \vdots \\ x_1 x_2 \dots x_d \end{pmatrix}$$

Inner-Product

Now the question is, we need to compute the $\phi(x)^T \phi(z)$: Kernel Function

$$\phi(x) = \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_d \\ x_1 x_2 \\ \vdots \\ x_{d-1} x_d \\ \vdots \\ x_1 x_2 \dots x_d \end{pmatrix}$$

$$\phi(x_i)^T \phi(x_j) = k(x_i, x_j)$$


The inner product $\phi(x)^T \phi(z)$ can be formulated as:

$$\phi(x)^T \phi(z) = 1 \cdot 1 + x_1 z_1 + x_2 z_2 + \dots + x_1 x_2 z_1 z_2 + \dots + x_1 \dots x_d z_1 \dots z_d = \prod_{k=1}^d (1 + x_k z_k)$$

The sum of 2^{2d} terms

The product of
 d terms

Inner-Product

- So with a finite training set of n samples, inner products are often pre-computed and stored in a Kernel Matrix

$$K_{ij} = \phi(x_i)^T \phi(x_j)$$

- If we store the matrix K , we only need to do simple inner-product look-ups
- Now instead of compute for w (which has the same dimensions with $\phi(x_i)$) we only need to compute α low-dimensional computations throughout the gradient descent algorithm. For prediction on the test data x_t

$$h(x_t) = \sum_{j=1}^n \alpha_j k(x_j, x_t)$$

Summary

- For kernel methods:
 - Re-express the model with inner-product
 - Replace inner-product terms with a kernel function
 - Compute the kernel term values and store them (basically, this learns interaction patterns among the training data)
 - Derive the new coefficients (e.g. α not w)
 - When testing, the prediction is based on the interaction of the training data and the new test data

We capture the interactions between data points in a very high dimensional space, and we never explicitly compute an instance in that space.