

# Applied Machine Learning Mini Project

Student Name: Aditya Sanjay Mhaske

Topic: Object Detection

## Problem Statement:

- The training set contains 100 categories with labels from 0 to 99. Each category has 5000 images. Each image is formatted as  $28 \times 28$  pixels. There are 500K images in total. You can download the dataset from <https://bit.ly/p556bonus>.
- The test set covers the same categories in the training set without any unseen categories. The test set has 100K images. We didn't provide the labels of the test set.

## 1. Methodology (Formulation of the problem)

Technologies used:

Keras, Tensorflow, and implemented CNN

The Convolutional Neural Network:

(CNN or ConvNet) is a subtype of Neural Network that is mainly used for applications in image recognition. Its built-in convolutional layer reduces the high dimensionality of images without losing their information. That is why CNN's are especially suited for this use case.

## 2. Data Preprocessing

Rescaled images to  $28 \times 28 \times 1$

```
'''
keras.layers.Rescaling(1./255, input_shape=(28, 28, 1)),
'''
```

## 3. Machine Learning model Implementation

```
In [20]: model = keras.Sequential([
keras.layers.Rescaling(1./255, input_shape=(28, 28, 1)),
keras.layers.Conv2D(32, (3,3), input_shape=(28,28,1), activation='relu'),
keras.layers.BatchNormalization(axis=-1),
keras.layers.Conv2D(32, (3,3), activation='relu'),
keras.layers.BatchNormalization(axis=-1),
keras.layers.MaxPooling2D(pool_size=(2,2)),
keras.layers.Conv2D(64, (3,3), activation='relu'),
keras.layers.BatchNormalization(axis=-1),
keras.layers.Conv2D(64, (3,3), activation='relu'),
keras.layers.BatchNormalization(axis=-1),
keras.layers.MaxPooling2D(pool_size=(2,2)),
keras.layers.Flatten(),
keras.layers.Dense(512),
keras.layers.BatchNormalization(),
keras.layers.Activation('relu'),
keras.layers.Dropout(0.2),
keras.layers.Dense(100),
keras.layers.Dense(100, activation=tf.nn.softmax)
])
```

# Model Summary :

Model: "sequential"		
Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 28, 28, 1)	0
conv2d (Conv2D)	(None, 26, 26, 32)	320
batch_normalization (Batch Normalization)	(None, 26, 26, 32)	128
conv2d_1 (Conv2D)	(None, 24, 24, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 24, 24, 32)	128
max_pooling2d (MaxPooling2D)	(None, 12, 12, 32)	0
conv2d_2 (Conv2D)	(None, 10, 10, 64)	18496
batch_normalization_2 (Batch Normalization)	(None, 10, 10, 64)	256
conv2d_3 (Conv2D)	(None, 8, 8, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 8, 8, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 512)	524800
batch_normalization_4 (Batch Normalization)	(None, 512)	2048
activation (Activation)	(None, 512)	0
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 100)	51300
dense_2 (Dense)	(None, 100)	10100
Total params: 654,008		
Trainable params: 652,600		
Non-trainable params: 1,408		

## 5. Results and Observations:

1. Training Accuracy: 74.2%
2. Validation accuracy: 72.2%

```
In [10]: model.fit(train_x,train_y,epochs=10)

Epoch 1/10
2022-12-04 21:09:37.562625: W tensorflow/core/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU frequency: 0
Hz
2022-12-04 21:09:38.417730: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:114] Plugin optimizer for device_type GPU is enabled.

14063/14063 [=====] - 251s 17ms/step - loss: 1.6619 - accuracy: 0.5759
Epoch 2/10
14063/14063 [=====] - 283s 20ms/step - loss: 1.2885 - accuracy: 0.6625
Epoch 3/10
14063/14063 [=====] - 241s 17ms/step - loss: 1.1890 - accuracy: 0.6864
Epoch 4/10
14063/14063 [=====] - 237s 17ms/step - loss: 1.1275 - accuracy: 0.7007
Epoch 5/10
14063/14063 [=====] - 237s 17ms/step - loss: 1.0820 - accuracy: 0.7115
Epoch 6/10
14063/14063 [=====] - 238s 17ms/step - loss: 1.0442 - accuracy: 0.7206
Epoch 7/10
14063/14063 [=====] - 236s 17ms/step - loss: 1.0133 - accuracy: 0.7272
Epoch 8/10
14063/14063 [=====] - 236s 17ms/step - loss: 0.9865 - accuracy: 0.7330
Epoch 9/10
14063/14063 [=====] - 239s 17ms/step - loss: 0.9628 - accuracy: 0.7379
Epoch 10/10
14063/14063 [=====] - 236s 17ms/step - loss: 0.9446 - accuracy: 0.7420

Out[10]: <keras.callbacks.History at 0x2bd9f2af0>

In [ ]:

In [11]: val_loss, val_acc = model.evaluate(test_x[:10000], test_y[:10000])
print('Validation accuracy:', val_acc)

2022-12-04 21:50:13.282074: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:114] Plugin optimizer for device_type GPU is enabled.

313/313 [=====] - 3s 8ms/step - loss: 1.0919 - accuracy: 0.7219
Validation accuracy: 0.7219000458717346
```

## 6. How to improve accuracy

By using a pre-trained model like resnet can be implemented to improve further accuracy. Else adding more layers to the neural network can further lead to improve accuracy.